

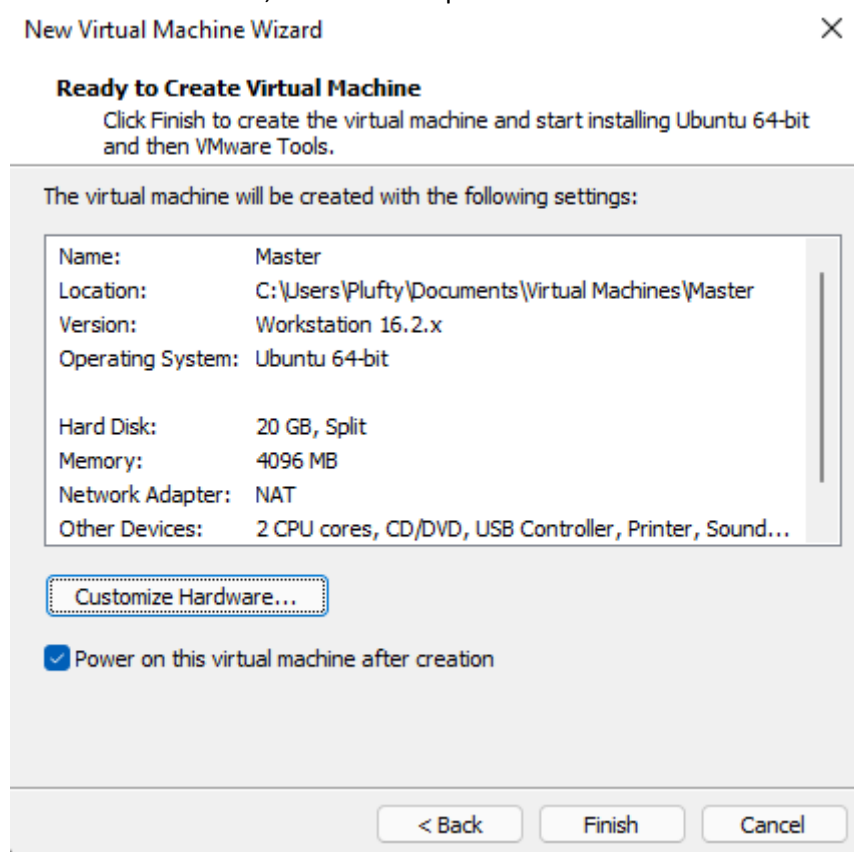
## Gleidson Vinícius Gomes Barbosa - 6331

### Cluster Kubernetes - Como fazer?

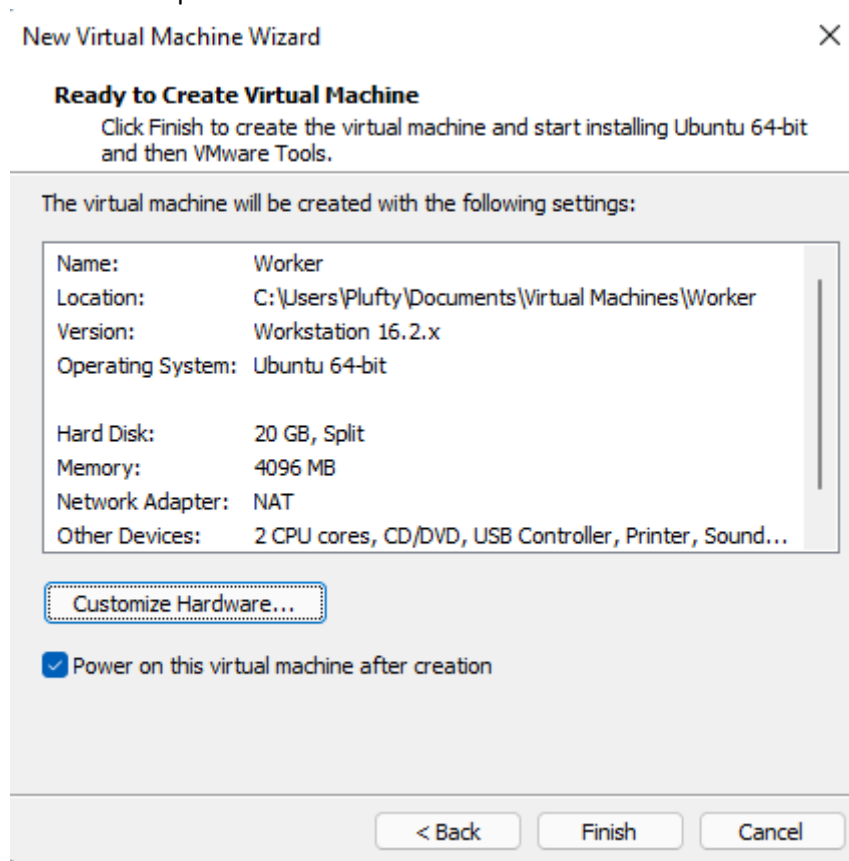
O procedimento será executado por mim em duas máquinas virtuais utilizando VMWare, então o primeiro passo será a criação de ambas as máquinas virtuais e ativar a virtualização das mesmas visto que será utilizado o KVM. Será instalado em ambas as máquinas o Ubuntu Server 20.04. Os procedimentos detalhados nesse relatório exceto pelos relacionados ao VMWare podem ser encontrados em:

- [Installing kubeadm | Kubernetes](#)
- [Install Docker Engine on Ubuntu | Docker Documentation](#)
- [Ports and Protocols | Kubernetes](#)
- [Como abrir portas no Ubuntu Server? \(vocepergunta.com\)](#)
- [kubernetes - Kubelet service is not running. It seems like the kubelet isn't running or healthy - Stack Overflow](#)
- [Creating a cluster with kubeadm | Kubernetes](#)
- [Run a Stateless Application Using a Deployment | Kubernetes](#)

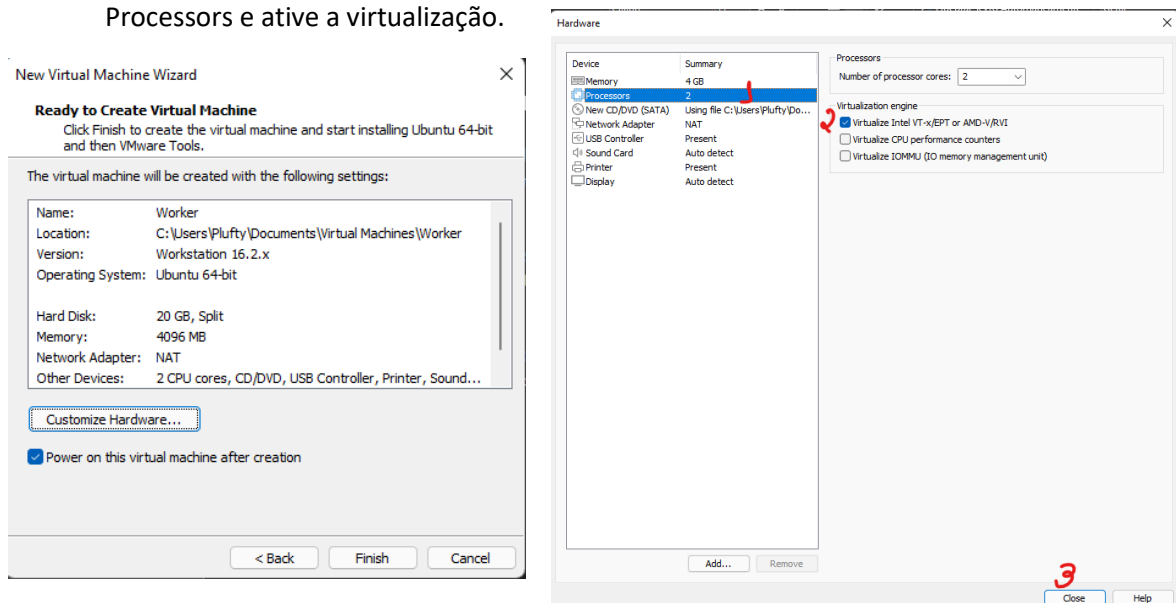
1. Criando a VM Master, ela será a máquina master desse cluster.




2. Criando a VM Worker que será a máquina cliente desse processo, ela será o único worker desse processo.



3. Ativando a virtualização no VMWare, ao clicar em Customize Hardware, selecione Processors e ative a virtualização.




#### 4. Instalação do SO, nesse caso irei utilizar o Ubuntu server 18.04




## A instalação do Docker e Kubernetes

Neste momento estaremos instalando o Docker em ambas as máquinas, por isso colocarei prints de apenas uma delas, basta repetir o procedimento em ambas.

1. Como estamos trabalhando com máquinas virtuais, estarei acessando primeiramente via SSH a máquina Master para trabalhar sem problemas de desempenho do VMWare.



2. Logo após, entrarei em modo root para não ter problemas com permissões. Para isso utiliza-se o comando sudo -i.



3. Atualizaremos os pacotes do linux com *apt update*

```
root@master:~# apt update
Hit:1 http://archive.ubuntu.com/ubuntu bionic InRelease
Hit:2 http://archive.ubuntu.com/ubuntu bionic-updates InRelease
Get:3 http://archive.ubuntu.com/ubuntu bionic-backports InRelease [74.6 kB]
Get:4 http://archive.ubuntu.com/ubuntu bionic/multiverse Sources [181 kB]
Get:5 http://security.ubuntu.com/ubuntu bionic-security InRelease [88.7 kB]
Get:6 http://archive.ubuntu.com/ubuntu bionic/restricted Sources [5,324 B]
Get:7 http://archive.ubuntu.com/ubuntu bionic/universe Sources [9,051 kB]
Get:8 http://security.ubuntu.com/ubuntu bionic-security/multiverse Sources [293 kB]
Get:9 http://security.ubuntu.com/ubuntu bionic-security/restricted Sources [22.7 kB]
Get:10 http://security.ubuntu.com/ubuntu bionic-security/main Sources [262 kB]
Get:11 http://security.ubuntu.com/ubuntu bionic-security/multiverse Sources [7,972 B]
Get:12 http://security.ubuntu.com/ubuntu bionic-security/main Translation-en [373 kB]
Get:13 http://security.ubuntu.com/ubuntu bionic-security/restricted amd64 Packages [646 kB]
Get:14 http://security.ubuntu.com/ubuntu bionic-security/restricted Translation-en [68.3 kB]
Get:15 http://security.ubuntu.com/ubuntu bionic-security/universe amd64 Packages [1,171 kB]
Get:16 http://security.ubuntu.com/ubuntu bionic-security/universe Translation-en [271 kB]
Get:17 http://security.ubuntu.com/ubuntu bionic-security/multiverse amd64 Packages [17.6 kB]
Get:18 http://security.ubuntu.com/ubuntu bionic-security/multiverse Translation-en [3,668 B]
Get:19 http://archive.ubuntu.com/ubuntu bionic/main Sources [829 kB]
Get:20 http://archive.ubuntu.com/ubuntu bionic/main Translation-en [516 kB]
Get:21 http://archive.ubuntu.com/ubuntu bionic/restricted amd64 Packages [2,184 B]
Get:22 http://archive.ubuntu.com/ubuntu bionic/restricted Translation-en [3,584 B]
Get:23 http://archive.ubuntu.com/ubuntu bionic/universe amd64 Packages [6,570 kB]
Get:24 http://archive.ubuntu.com/ubuntu bionic/universe Translation-en [5,941 kB]
Get:25 http://archive.ubuntu.com/ubuntu bionic/multiverse amd64 Packages [151 kB]
Get:26 http://archive.ubuntu.com/ubuntu bionic/multiverse Translation-en [108 kB]
Get:27 http://archive.ubuntu.com/ubuntu bionic-updates/restricted Sources [15.5 kB]
Get:28 http://archive.ubuntu.com/ubuntu bionic-updates/multiverse Sources [15.9 kB]
Get:29 http://archive.ubuntu.com/ubuntu bionic-updates/main Sources [522 kB]
Get:30 http://archive.ubuntu.com/ubuntu bionic-updates/universe Sources [470 kB]
Get:31 http://archive.ubuntu.com/ubuntu bionic-updates/main Translation-en [465 kB]
Get:32 http://archive.ubuntu.com/ubuntu bionic-updates/restricted amd64 Packages [671 kB]
Get:33 http://archive.ubuntu.com/ubuntu bionic-updates/restricted Translation-en [92.1 kB]
Get:34 http://archive.ubuntu.com/ubuntu bionic-updates/universe amd64 Packages [1,793 kB]
Get:35 http://archive.ubuntu.com/ubuntu bionic-updates/universe Translation-en [388 kB]
Get:36 http://archive.ubuntu.com/ubuntu bionic-updates/multiverse amd64 Packages [24.2 kB]
Get:37 http://archive.ubuntu.com/ubuntu bionic-updates/multiverse Translation-en [5,980 B]
Get:38 http://archive.ubuntu.com/ubuntu bionic-backports/universe Sources [6,600 B]
Get:39 http://archive.ubuntu.com/ubuntu bionic-backports/main Sources [5,476 B]
Get:40 http://archive.ubuntu.com/ubuntu bionic-backports/main amd64 Packages [10.3 kB]
Get:41 http://archive.ubuntu.com/ubuntu bionic-backports/main Translation-en [4,824 B]
Get:42 http://archive.ubuntu.com/ubuntu bionic-backports/universe amd64 Packages [11.3 kB]
Get:43 http://archive.ubuntu.com/ubuntu bionic-backports/universe Translation-en [5,772 B]
Fetched 32.2 MB in 1min 43s (311 kB/s)
Reading package lists... Done
Building dependency tree
Reading state information... Done
27 packages can be upgraded. Run 'apt list --upgradable' to see them.
```

4. O primeiro passo que faremos é desligar a memória swap com o comando *swapoff -a*

```
root@master:~# swapoff -a
```

5. Depois instalaremos o Docker com o comando *curl -fsSL https://get.docker.com | bash* após completar a instalação conferimos a versão com *docker --version*

```
root@master:~# curl -fsSL https://get.docker.com | bash
# Executing docker install script, commit: 93d2499759296ac1f9c510605fef85052a2c32be
# sh -c 'set -e; apt-get update --no-install-recommends && { docker pull --quiet docker:latest; }'
root@master:~# docker --version
Docker version 20.10.13, build a224086
```

6. Nesse ponto, utilizaremos uma série de comandos para carregar os módulos necessários para o kubernetes e configurar as bridges, seguem os comandos:

```
modprobe br_netfilter
lsmod | grep br_netfilter
cat <<EOF | sudo tee /etc/modules-load.d/k8s.conf
br_netfilter
EOF
cat <<EOF | sudo tee /etc/sysctl.d/k8s.conf
net.bridge.bridge-nf-call-ip6tables = 1
net.bridge.bridge-nf-call-iptables = 1
EOF
sudo sysctl --system
```

```
root@master:~# modprobe br_netfilter
root@master:~# lsmod | grep br_netfilter
br_netfilter      24576  0
bridge            155648  1 br_netfilter
root@master:~# cat <<EOF | sudo tee /etc/modules-load.d/k8s.conf
> br_netfilter
> EOF
cat <<EOF | sudo tee /etc/sysctl.d/k8s.conf
net.bridge.bridge-nf-call-ip6tables = 1
net.bridge.bridge-nf-call-iptables = 1
EOF
sudo sysctl --system
```

7. Agora verificaremos se as portas requeridas pelo kubernetes estão abertas. Para abrir usaremos os seguintes comandos:

```
iptables -I INPUT -p tcp --dport 6443 -j ACCEPT
```

```
iptables -I INPUT -p tcp --dport 2379 -j ACCEPT
iptables -I INPUT -p tcp --dport 2389 -j ACCEPT
iptables -I INPUT -p tcp --dport 10250 -j ACCEPT
iptables -I INPUT -p tcp --dport 10259 -j ACCEPT
iptables -I INPUT -p tcp --dport 10257 -j ACCEPT
```

8. Agora atualizaremos os pacotes do Linux com o comando apt-get update

```
root@master:~# apt-get update
Get:1 http://security.ubuntu.com/ubuntu bionic-security InRelease [88.7 kB]
Hit:2 http://archive.ubuntu.com/ubuntu bionic InRelease
Get:3 http://archive.ubuntu.com/ubuntu bionic-updates InRelease [88.7 kB]
Get:4 http://archive.ubuntu.com/ubuntu bionic-backports InRelease [74.6 kB]
Hit:5 https://download.docker.com/linux/ubuntu bionic InRelease
Fetched 252 kB in 3s (76.7 kB/s)
Reading package lists... Done
```

9. Feito isso, alguns pré requisitos do kubernetes, como os certificados com o comando apt-get install -y apt-transport-https ca-certificates curl

```
root@master:~# sudo apt-get install -y apt-transport-https ca-certificates curl
Reading package lists... Done
Building dependency tree
Reading state information... Done
ca-certificates is already the newest version (20210119~18.04.2).
curl is already the newest version (7.58.0-2ubuntu3.16).
apt-transport-https is already the newest version (1.6.14).
0 upgraded, 0 newly installed, 0 to remove and 280 not upgraded.
```

10. Adicionaremos uma chave publica do google cloud com o comando:

```
curl -fsSL /usr/share/keyrings/kubernetes-archive-keyring.gpg
https://packages.cloud.google.com/apt/doc/apt-key.gpg
```

logo após adicionaremos o repositório kubernetes com o comando:

```
echo "deb [signed-by=/usr/share/keyrings/kubernetes-archive-keyring.gpg]
https://apt.kubernetes.io/ kubernetes-xenial main" | sudo tee
/etc/apt/sources.list.d/kubernetes.list
```

```
root@master:~# curl -fsSL /usr/share/keyrings/kubernetes-archive-keyring.gpg https://packages.cloud.google.com/apt/doc/apt-key.gpg
root@master:~# echo "deb [signed-by=/usr/share/keyrings/kubernetes-archive-keyring.gpg] https://apt.kubernetes.io/ kubernetes-xenial main" | sudo tee /etc/apt/sources.list.d/kubernetes.list
deb [signed-by=/usr/share/keyrings/kubernetes-archive-keyring.gpg] https://apt.kubernetes.io/ kubernetes-xenial main
root@master:~#
```

11. Após adicionar o repositório, atualizaremos novamente os pacotes e por fim instalaremos o kubernetes com a seguinte sequência de comandos:

```
sudo apt-get update
sudo apt-get install -y kubelet kubeadm kubectl
sudo apt-mark hold kubelet kubeadm kubectl
root@master:~# sudo apt-get update
root@master:~# sudo apt-get install -y kubelet kubeadm kubectl
root@master:~# sudo apt-mark hold kubelet kubeadm kubectl
```

12. Executaremos o comando kubeadm --init para iniciar nosso mastere ao fim deste comando será gerado um token, lembre-se de salvá-lo, ele será necessário no futuro. No meu caso é o seguinte:

```
kubeadm join 192.168.254.147:6443 --token vmurq9.m20zvlq7pd7myruf \
discovery-token-ca-cert-hash
```

sha256:cd930035a44414aba217c3081632ad37035caae039032dffd41b60bde26e0d

```
Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config

Alternatively, if you are the root user, you can run:

export KUBECONFIG=/etc/kubernetes/admin.conf

You should now deploy a pod network to the cluster.
Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
https://kubernetes.io/docs/concepts/cluster-administration/addons/

Then you can join any number of worker nodes by running the following on each as root:

kubeadm join 192.168.254.147:6443 --token vmurg9.m20zvlq7pd7myruf \
--discovery-token-ca-cert-hash sha256:cd930035a44414aba217c3081632ad37035caae039032dffd41b60bde26e0d
root@master:~#
```

13. Agora iremos configurar o kubernetes de acordo com a própria resposta anterior utilizando os comandos indicados:

`mkdir -p $HOME/.kube`

`sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config`

`sudo chown $(id -u):$(id -g) $HOME/.kube/config`

```
root@master:~# mkdir -p $HOME/.kube
root@master:~# sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
root@master:~# sudo chown $(id -u):$(id -g) $HOME/.kube/config
root@master:~# kubectl getnodes
error: unknown command "getnodes" for "kubectl"
```

14. Verificamos com o comando `kubectl get nodes` que já temos nosso node Master no estado de NotReady.

```
root@master:~# kubectl get nodes
NAME          STATUS    ROLES          AGE   VERSION
master        NotReady  control-plane, 64m   v1.23.4
```

15. Agora iremos adicionar nosso woker no nosso cluster. Precisaremos do comando com o token no passo 12, basta copiá-lo e executar na máquina worker.

```
root@worker:~# kubeadm join 192.168.254.147:6443 --token vmurg9.m20zvlq7pd7myruf \
--discovery-token-ca-cert-hash sha256:cd930035a44414aba217c3081632ad37035caae039032dffd41b60bde26e0d
[preflight] Running pre-flight checks
[preflight] Reading configuration from the cluster...
[preflight] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -o yaml'
W0311 04:10:52.793496 72410 utils.go:69] The recommended value for "resolvConf" in "KubeletConfiguration" is: /run/
systemd/resolve/resolv.conf; the provided value is: /run/systemd/resolve/resolv.conf
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Starting the kubelet
[kubelet-start] Waiting for the kubelet to perform the TLS Bootstrap...

This node has joined the cluster:
 * Certificate signing request was sent to apiservert and a response was received.
 * The Kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the control-plane to see this node join the cluster.
```

16. Agora verificaremos no nosso master com o comando `kubectl get nodes` que temos o master e o worker no nosso cluster, ambos em estado NotReady.

```
root@master:~# kubectl get nodes
NAME          STATUS    ROLES          AGE   VERSION
master        NotReady  control-plane, 74m   v1.23.4
worker        NotReady  <none>         14s   v1.23.4
```

17. Agora executaremos o comando que irá deixar nosso status como pronto:

`kubectl apply -f "https://cloud.weave.works/k8s/net?k8s-version=$(kubectl version | base64 | tr -d '\n')"`

```
root@master:~# kubectl apply -f "https://cloud.weave.works/k8s/net?k8s-version=$(kubectl version | base64 | tr -d '\n')
serviceaccount/weave-net created
clusterrole.rbac.authorization.k8s.io/weave-net created
clusterrolebinding.rbac.authorization.k8s.io/weave-net created
role.rbac.authorization.k8s.io/weave-net created
rolebinding.rbac.authorization.k8s.io/weave-net created
daemonset.apps/weave-net created
```

18. Executamos novamente o comando `kubectl get nodes` e observamos que nosso cluster está pronto.

```
root@master:~# kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
master	Ready	control-plane,master	89m	v1.23.4
worker	Ready	<none>	15m	v1.23.4

## Fazendo o Deployment de uma aplicação

Agora faremos um deployment de uma aplicação no nosso cluster kubernetes. Para esse passo escolhemos a aplicação nginx.

1. Para isso iremos iniciar criando um arquivo yaml com o comando `kubectl apply -f https://k8s.io/examples/application/deployment.yaml`

```
root@master:~# kubectl apply -f https://k8s.io/examples/application/deployment.yaml
deployment.apps/nginx-deployment created
```

2. Após isso iremos verificar as informações sobre nosso deployment com o comando `kubectl describe deployment nginx-deployment`

```
root@master:~# kubectl describe deployment nginx-deployment
```

Name: nginx-deployment  
Namespace: default  
CreationTimestamp: Sat, 12 Mar 2022 23:47:06 +0000  
Labels: <none>  
Annotations: deployment.kubernetes.io/revision: 1  
Selector: app=nginx  
Replicas: 2 desired | 2 updated | 2 total | 2 available | 0 unavailable  
StrategyType: RollingUpdate  
MinReadySeconds: 0  
RollingUpdateStrategy: 25% max unavailable, 25% max surge

Pod Template:

Labels: app=nginx  
Containers:

nginx:

Image: nginx:1.14.2  
Port: 80/TCP  
Host Port: 0/TCP  
Environment: <none>  
Mounts: <none>  
Volumes: <none>

Conditions:

Type	Status	Reason
Available	True	MinimumReplicasAvailable
Progressing	True	NewReplicaSetAvailable

OldReplicaSets: <none>  
NewReplicaSet: nginx-deployment-9456bbbf9 (2/2 replicas created)

Events:

Type	Reason	Age	From	Message
Normal	ScalingReplicaSet	3m20s	deployment-controller	Scaled up replica set nginx-deployment-9456bbbf9 to 2

3. Depois verificaremos os pods gerados com o comando `kubectl get pods -l app=nginx`

```
root@master:~# kubectl get pods -l app=nginx
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-deployment-9456bbbf9-7hqdc	1/1	Running	0	4m26s
nginx-deployment-9456bbbf9-z5st5	1/1	Running	0	4m26s

4. Agora podemos verificar as informações sobre os pods gerados com o comando `kubectl describe pod <pod-name>` no nosso caso, usaremos como exemplo o pod de nome nginx-deployment-9456bbbf9-7hqdc



```

root@master:~# kubectl describe pod nginx-deployment-9456bbbf9-7hqdc
Name:          nginx-deployment-9456bbbf9-7hqdc
Namespace:     default
Priority:       0
Node:          worker/192.168.254.148
Start Time:    Sat, 12 Mar 2022 23:47:06 +0000
Labels:        app=nginx
               pod-template-hash=9456bbbf9
Annotations:   <none>
Status:        Running
IP:            10.44.0.2
IPs:
  IP:          10.44.0.2
Controlled By: ReplicaSet/nginx-deployment-9456bbbf9
Containers:
  nginx:
    Container ID:  docker://0fcc1cc6bce384f33ed1bfa74b000697f83ae5280abdc9842dbdff5b01753f8
    Image:         nginx:1.14.2
    Image ID:      docker-pullable://nginx@sha256:f7988fb6c02e0ce69257d9bd9cf37ae20a60f1df7563c3a2a6abe24160306b8d
    Port:         80/TCP
    Host Port:     0/TCP
    State:        Running
      Started:     Sat, 12 Mar 2022 23:47:19 +0000
      Ready:       True
      Restart Count: 0
    Environment:   <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-twbp8 (ro)
Conditions:
  Type              Status
  Initialized        True
  Ready             True
  ContainersReady    True
  PodScheduled       True
Volumes:
  kube-api-access-twbp8:
    Type:              Projected (a volume that contains injected data from multiple sources)
    TokenExpirationSeconds: 3607
    ConfigMapName:      kube-root-ca.crt
    ConfigMapOptional:  <nil>
    DownwardAPI:        true
  QoS Class:          BestEffort
  Node-Selectors:     <none>
  Tolerations:        node.kubernetes.io/not-ready:NoExecute op=Exists for 300s
                     node.kubernetes.io/unreachable:NoExecute op=Exists for 300s
Events:
  Type    Reason      Age   From          Message
  ----    -
  Normal  Scheduled   5m29s  default-scheduler  Successfully assigned default/nginx-deployment-9456bbbf9-7hqdc to worker
  Normal  Pulling     5m28s  kubelet        Pulling image "nginx:1.14.2"
  Normal  Pulled      5m17s  kubelet        Successfully pulled image "nginx:1.14.2" in 11.351030303s
  Normal  Created     5m16s  kubelet        Created container nginx
  Normal  Started     5m16s  kubelet        Started container nginx

```

Verificamos assim que nosso pod está rodando em nosso worker, ou seja, nosso deployment funcionou, Uhu!

## Erro no Kubelet

Fique atento a esse tópico pois temos alguns passos que só serão executados no nosso master, esses passos estão marcados com a tag **[APENAS NA MÁQUINA MASTER]**, se por acaso for executado no worker, pode comprometer todo nosso cluster.



1. Existe a possibilidade de alguns erros serem apresentados devido a problemas com o kubelet, o seguinte erro é apresentado:

```
Unfortunately, an error has occurred:
    timed out waiting for the condition

This error is likely caused by:
- The kubelet is not running
- The kubelet is unhealthy due to a misconfiguration of the node in some way (required cgroups disabled)

If you are on a systemd-powered system, you can try to troubleshoot the error with the following commands:
- 'systemctl status kubelet'
- 'journalctl -xeu kubelet'

Additionally, a control plane component may have crashed or exited when started by the container runtime.
To troubleshoot, list all containers using your preferred container runtimes CLI.

Here is an example how you may list all Kubernetes containers running in docker:
- 'docker ps -a | grep kube | grep -v pause'
Once you have found the failing container, you can inspect its logs with:
- 'docker logs CONTAINERID'
```

error execution phase wait-control-plane: couldn't initialize a Kubernetes cluster

To see the stack trace of this error execute with --v=5 or higher

2. Nesse caso seguiremos os seguintes procedimentos:  
Primeiro desabilitaremos novamente o swap e após isso verificaremos o grupo do docker com a seguinte sequência de comandos:

```
sudo sed -i 's/^/#/' /etc/fstab
```

sudo swapoff -a

`docker info |grep -i cgroup`

```
root@master:~# sudo sed -i 's/swap / s/^/#/' /etc/fstab
root@master:~# sudo swapoff -a
root@master:~# docker info |grep -i cgroup
```

3. Após isso editaremos o arquivo com o comando `nano /etc/docker/daemon.json` e adicionaremos o seguinte texto:

```
{
  "exec-opts": ["native.cgroupdriver=systemd"]
}
```

```
root@master:~# nano /etc/docker/daemon.json
```

```
GNU nano 2.9.3
{
  "exec-opts": ["native.cgroupdriver=systemd"]
}
```

4. Agora executaremos a seguinte sequência de comandos para reiniciar alguns serviços e verificar o kubelet:

```
sudo systemctl daemon-reload
```

```
sudo systemctl restart docker
```

```
sudo systemctl restart kubelet
```

[illegible]

5. **[APENAS NA MÁQUINA MASTER]** Depois iremos refazer a fase de tokens do kubeadm com o comando

```
kubeadm init phase bootstrap-token
```

```

root@master:~# kubeadm init phase bootstrap-token
[bootstrap-token] using token: bdeufs.pqteslyt4dykej
[bootstrap-token] Configuring bootstrap tokens, cluster-info ConfigMap, RBAC Roles
[bootstrap-token] configured RBAC rules to allow Node Bootstrap tokens to get nodes
[bootstrap-token] configured RBAC rules to allow Node Bootstrap tokens to post CSRs in order for nodes to get long term certificate credentials
[bootstrap-token] configured RBAC rules to allow the csrapprover controller automatically approve CSRs from a Node Bootstrap token
[bootstrap-token] configured RBAC rules to allow certificate rotation for all node client certificates in the cluster
[bootstrap-token] Creating the "cluster-info" ConfigMap in the "kube-public" namespace

```

6. Agora desabilitaremos o firewall com o comando *ufw disable*

```

root@master:~# ufw disable
Firewall stopped and disabled on system startup

```

7. **[APENAS NA MÁQUINA MASTER]** Finalmente executaremos o comando *kubeadm*  
*init --ignore-preflight-errors=all*

```

root@master:~# kubeadm init --ignore-preflight-errors=all

```

Corrigimos nossos erros, agora voltamos ao passo que paramos e seguimos para concluir nosso cluster, Felicidade meus amigos!

